Introduction to LangChain

LangChain is a framework for building applications powered by large language models (LLMs). It provides modular components such as prompts, models, retrievers, chains and agents.

Document loaders are one of the core building blocks. They read data from many sources (files, directories, web pages, databases) and convert it into Document objects containing page_content and metadata.

Introduction to LangChain (continued)

A typical RAG pipeline in LangChain has these steps: load documents, split them into chunks, embed the chunks, store them in a vector store, and retrieve relevant chunks at query time.

Text splitters such as RecursiveCharacterTextSplitter help keep chunks at a manageable size while preserving context.