Vector Databases Overview

A vector database stores embeddings, which are numerical representations of text, images or other data. Similar items have vectors that are close together in the embedding space.

Popular vector stores used with LangChain include FAISS, Chroma, Pinecone and Weaviate.

Vector Databases Overview (continued)

Similarity search finds the nearest vectors to a query vector using metrics like cosine similarity or Euclidean distance.

Vector databases are the backbone of semantic search and retrieval-augmented generation (RAG) systems.